

Memory management in Python - how it works and is it worth messing with it?

Is it possible to have a memory leak in Python? What is the performance cost of garbage collector? When can I run into problems with memory and how can I solve them?

Maciej Pytel



02/03/2015

About me

Software developer @ CodiLime. I code (mainly) in Python. I like to know how stuff works "inside".

Disclaimer

Which Python?

This talk is about Python 2.7.
By "Python" I mean CPython.

Examples

<https://github.com/MaciekPytel/python-memory-talk>

- 1 Just a bit of theory
- 2 Why is it worth knowing this stuff? What (and how) can we do?
 - Example: RESTful server
 - Memory leaks
- 3 A few useful tools
- 4 More theory, or how it all works?

Section 1

Just a bit of theory

Reference counting

- Every object has a reference counter.
- When the counter reaches 0 the object is *immediately* deallocated.
- What about reference cycles?

Garbage collector

Mission: delete objects in unreachabeable reference cycles.

Solution: garbage collector.

- Automatically run by Python interpreter every once in a while.
- Doesn't explicitly deallocate any objects - just breaks reference cycles.
- Freezes the process for garbage collection.
- Expensive: cost is (at least) linear with regard to the number of references in program.

Weak generational hypothesis

Most objects live either for a very short or a very long time.

Can this help us optimise GC?

Generational GC

- Divide objects into 3 generations (0, 1, 2).
- Each newly allocated object is added to generation 0.
- Long living objects are promoted to higher generations.
- Generation 0 is made of relatively few objects with relatively large chance of being ready for deallocation (according to weak generational hypothesis).

- n -th generation garbage collection analyses objects of generation $0..n$.
- Any object that survived n -th generation collection is promoted to generation $n + 1$ (technically $\min(2, n + 1)$).
- 0-th generation collection is cheap and done often.
- Higher generation collections are more expensive and performed much less often.

Weak references

- Weak reference is a reference which doesn't increase reference counter of the object it points to.
- Existence of weak reference doesn't prevent deallocation of its target.
- In Python provided by *weakref* module.
- Dereferencing weak reference to deallocated object yields None.

Section 2

Why is it worth knowing this stuff? What (and how) can we do?

Subsection 1

Example: RESTful server

Experiment

- Server:
 - Flask-RESTful + gunicorn.
 - Domain made of 2 classes: Task and Subtask.
 - Natural one-to-many relationship (Each Task has multiple Subtasks).
 - Simple API supporting creation, retrieval and deletion of Tasks and Subtasks.
- Test: run a bunch of clients sending requests in a loop and measure server performance.

Task


```
class Task(object):  
    def __init__(self, name):  
        self.name = name  
        self._subtasks = {}  
  
    def add_subtask(self, name):  
        self._subtasks[name] = Subtask(name, self)  
  
    def get_subtask(self, name):  
        return self._subtasks.get(name, None)
```

Subtask version 1

```
class Subtask(object):  
    def __init__(self, name, task):  
        self.name = name  
        self._task = task  
  
    def get_task(self):  
        return self._task
```


- **GET** `/tasks/{task_id}/` - return all Subtasks corresponding to a given Task
- **PUT** `/tasks/{task_id}/` - create a Task with a given id
- **DELETE** `/tasks/{task_id}/` - delete a Task and all corresponding Subtasks
- **GET** `/subtasks/{subtask_id}/` - return a Subtask and its parent Task
- **PUT** `/subtasks/{task_id}/{subtask_id}` - create a Task with a given id and assign it to a Task

Results




Subtask version 2

```
import weakref


class Subtask(object):
    def __init__(self, name, task):
        self.name = name
        self._task = weakref.ref(task)

    def get_task(self):
        return self._task()
```


Results?



Results



Results



What have we learned?

- Relying on GC is expensive.
- Specifically 2nd generation GC takes a long time (2-2.5s in example server).
- Lower generation collections are significantly cheaper (no noticeable impact on request time for client).
- 2nd generation collection is run very rarely.
- Python is able to reuse freed memory, but it doesn't seem to return it to OS.

Freeing memory

- Python uses its own memory allocator (we'll skip the details).
- Allocator can return the memory to OS^{*}, but:
 - Memory is allocated in blocks and only released once all objects in a block are deallocated (fragmentation!).
 - Many objects are returned to pool for reuse (ex. tuples).
 - Memory allocated for integers will never be freed or used for anything else.
 - Same with floats.
 - ...

Conclusion

If our program is having memory problems we may:

- Remove references to objects as soon as we don't need them anymore.
- Avoid *unnecessary* reference cycles.
- In most cases when they are necessary we can use weakrefs.
- Or explicitly break the cycle when we don't need the object anymore (ex. by defining and calling *delete* method).
- *xrange* instead of *range* in long loops makes a difference. Seriously, use generators.

Subsection 2

Memory leaks

Memory leaks

There are a few possible reasons for leaks:

- 1 Memory leak in c/c++ module.
- 2 Memory allocated for integers or floats will never be freed or used for anything else.
- 3 "Forgotten" references.
- 4 Garbage collector won't free *any* object in a cycle where at least a single object defines a finalizer (`__del__`).

"Forgotten" references

There are a few places where unexpected references can hide and keep your objects alive:

- 1 `sys.exc_info()` returns information about last exception handled in this stack frame. Part of this information is a traceback object containing whole stack state at the time of exception. If you catch an exception in your main loop: Oops.
- 2 Closures, `functools.partial`, etc.

Finalizer (___del___)

- If an object defines `___del___` method it will be called just before object deallocation.
- Object can create a reference to itself in `___del___` - this will prevent deallocation.
- **GC will never deallocate an object defining `___del___`!**
- In other words the whole cycle containing such an object will never be broken.
- **Warning:** `___del___` is used in many existing libraries (including Python standard library!).

As we can see finalizers can cause problems. How do we deal with them?

- Don't use finalizers. Context manager ("with") is almost always a better solution.
- If you really need a finalizer - make sure it's never a part of reference cycle (think decomposition).
- As a last resort: `weakref.ref` takes an optional callback parameter. This callback will be called after the object is deallocated. Problem: once the callback gets invoked the object no longer exists, so we need to keep the necessary state externally.

weakref based finalizer

```
class FileWrapper(object):
    _weakrefs = set()

    @classmethod
    def _delegated_close(cls, file_object, w):
        file_object.close()
        cls._weakrefs.remove(w)

    def __init__(self, name, mode):
        self._f = open(name, mode)
        self._weakrefs.add(weakref.ref(
            self,
            functools.partial(self._delegated_close,
                               self._f)))
```

Section 3

A few useful tools


Debugger, part of Python standard library. Allows you to inspect the state of your program, execute code and evaluate expressions in any stack frame. The basic Python debugging tool.

Python standard library module.

- **gc.collect(generation=2)** - trigger GC manually.
- **gc.enable()** / **gc.disable()** - enable/disable automatic GC.
- **gc.garbage** - a list of all objects that should be deleted, but define `__del__` method and are a part of reference cycle.
- You can set debugging flags, ex. to print GC statistics with variable level of details.
- Can give you a full list of objects referring a given object.

Simple library for analysys and visulisation of reference graph. Very useful when looking for memory leaks.

- Available via pip.
- Visualise reference graphs.
- Count objects in memory by type.
- Also display change in number of objects of given type in time.
- Sufficient to diagnose 99% of memory problems.



- Memory profiler.
- Incredible number of available functions: analyse paths in reference graph, reference graph spanning trees (!), grouping objects by various equality relationships (ex. class, module the originates from), ...
- Provides hooks for tracking objects defined in C module.
- Steep learning curve, not the best documentation and lack of good tutorials.
- For this 1% of situations when objgraph simply isn't enough :)

Section 4

More theory, or how it all works?


Generations

- Each generation is a doubly-linked list.
- Each newly allocated object is added to generation 0.
- Before n-th generation GC lower generations' lists are merged to n-th generation list.
- After n-th generation GC is done surviving objects' list is merged to generation $n + 1$.


GC algorithm

- 1 (For generation 1 and 2) Merge lower generation into current generation (doubly-linked list merge).
- 2 Iterate over objects on the list. Check all outgoing references (non-recursively) - if they point to other objects on the list decrease target's refcount.
- 3 Iterate over objects on the list. Move any object with refcount equal to 0 to a new *unreachable* list.
- 4 Iterate over objects remaining on the original list. Check all outgoing references and move any target on *unreachable* list back to the original list.
- 5 Restore original refcounts on all objects.
- 6 Merge original list to generation $n + 1$.
- 7 For each object on *unreachable* list remove all references originating from that object.

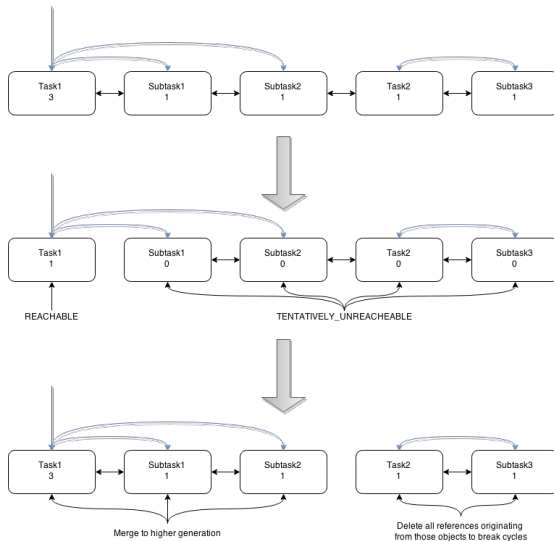
Example



Example



GC - visualisation



- The above is the idea of the algorithm. The real implementation is a bit more complex.
- Details in Python source `/Modules/gcmodule.c` - very well commented.

Programmers waste enormous amounts of time thinking about the speed of noncritical parts of their programs (...). We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.

- Donald Knuth